



Exceptional service in the national interest

# The Structural Simulation Toolkit (SST)

HPCA 2024

**Presented by: Patrick Lavin, Sandia National Laboratories**

**Clay Hughes, Sandia National Laboratories**

**Content by: SST Team**

Edinburgh, United Kingdom

Sunday, March 3

SAND2024-02470C

Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia LLC, a wholly owned subsidiary of Honeywell International Inc. for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.





# Welcome!

## Part 1: Introduction to SST

**8:00 – 10:00**

SST Overview

Basic Simulation in SST

A Tour of SST Elements

Break

*10:00 - 10:20*

Part 2: Node-Level Simulation

*10:20 – 11:00*

Part 3: System-Scale Simulation

*11:00 - 12:20*



# Instructors



Patrick Lavin

[prlavin@sandia.gov](mailto:prlavin@sandia.gov)



Clay Hughes

[chughes@sandia.gov](mailto:chughes@sandia.gov)



Scott Hemmert

[kshemme@ncsu.edu](mailto:kshemme@ncsu.edu)



## Tutorial Container

Download the container for this tutorial:

<https://github.com/sstsimulator/sst-tutorials/releases/tag/hpca2024>

Slides & container builder:

<https://github.com/sstsimulator/sst-tutorials/tree/master/hpca2024>

Using the container:

```
apptainer run --writable-tmpfs sst.sif
```



# References

## Websites

- Information on installing SST, and links to everything else you see here
  - <http://www.sst-simulator.org/>
- Documentation on SST's key interfaces, overview of all the elements, and more!
  - <http://sst-simulator.org/sst-docs/>
- Code
  - <https://github.com/sstsimulator>

## Important Pages

- Configuration File Format: <http://sst-simulator.org/SSTPages/SSTUserPythonFileFormat/>
- Doxygen Documentation: [http://sst-simulator.org/SSTDoxygen/13.0.0\\_docs/html/](http://sst-simulator.org/SSTDoxygen/13.0.0_docs/html/)
- Developer FAQ: <http://sst-simulator.org/SSTPages/SSTTopDocDeveloperInfo/>
- Building SST
  - Quick start: [http://sst-simulator.org/SSTPages/SSTBuildAndInstall\\_13dot1dot0\\_SeriesQuickStart/](http://sst-simulator.org/SSTPages/SSTBuildAndInstall_13dot1dot0_SeriesQuickStart/)
  - Detailed: [http://sst-simulator.org/SSTPages/SSTBuildAndInstall\\_13dot1dot0\\_SeriesDetailedBuildInstructions/](http://sst-simulator.org/SSTPages/SSTBuildAndInstall_13dot1dot0_SeriesDetailedBuildInstructions/)



## Learning Objectives

This section of the tutorial will cover the following topics:

1. The basic structure of the SST project
2. How to build a simulation in SST with existing components
3. How to find information about components with sst-info
4. How to gather information from your simulation with statistics
5. A summary of the many available components
6. Where to get more info and help





## Motivation: Interoperability

There is a rich selection of open-source simulators

But not a solid ecosystem for modeling systems

- Tightly-entangled components make modifications complex
  - E.g., assumptions about caching or address mapping are pervasive
- Most simulator integrations are ad-hoc, not lasting
- Significant performance problems with tying many simulators together

Wants:

- Enable “mix-and-match” of existing models to create custom systems
- Encourage disentangled models with clean interfaces for swapping functionality
  - Bricks not buildings
- Low effort, high performance parallel simulation
- Continuous path from low-fidelity/fast modeling to high-fidelity/slow models





# The Structural Simulation Toolkit

## Goals

- Create a standard architectural *simulation framework* for HPC\*
- Ability to evaluate future systems on DOE/DOD workloads
- *Use supercomputers to design supercomputers*

## Technical Approach

- **Parallel** Discrete Event core
  - With conservative optimization over MPI/Threads
- **Interoperability**
  - Node and system-scale models
- **Multi-scale**
  - Detailed (~cycle) and simple models that interoperate
- **Open**
  - Open Core, non-viral, modular

## Status

- Parallel framework (*SST Core*)
- Integrated libraries of components (*Elements*)
- Current Release (13.1.0)
  - Two releases per year

## Consortium

- “Best of Breed” simulation suite
- Used in industry, academia, and at labs







# The SST Approach

## Parallel Discrete-Event Simulator Framework (*SST Core*)

- Flexible framework enables multitude of custom “simulators”
- Demonstrated scaling to over 512 processors running a million+ components

## Comes with many built-in simulation models (*SST Elements*)

- Processors, memories, networks


## Open API

- Easily extensible with new models
- Modular framework
- Open-source core

## Time-scale independent core

- Handles Micro-, Meso-, Macro-scale simulations

C++, Python





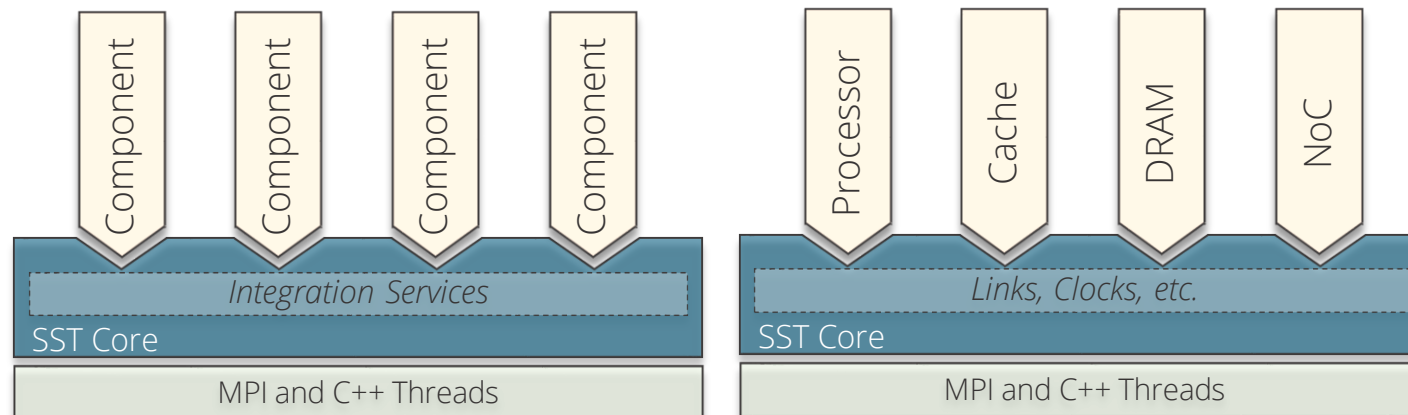
# SST Architecture

## SST **Core** framework

- The backbone of simulation
- Provides utilities and interfaces for simulation components (models)
  - Clocks, event exchange, statistics and parameter management, parallelism support, etc.

## SST **Element** libraries

- Libraries of components that perform the actual simulation
- Elements include processors, memory, network, etc.
  - Includes many existing simulators: DRAMSim2, Spike, HMCsim, Ramulator, etc.



The graphic features a central dark blue diamond with the text "High-level View" in white. This diamond is surrounded by a white border and is flanked by two diagonal lines of colorful segments (cyan, orange, green, red, purple) that extend towards the corners of the frame. The background is white with faint, light blue abstract shapes.

# High-level View

# Getting and Installing SST

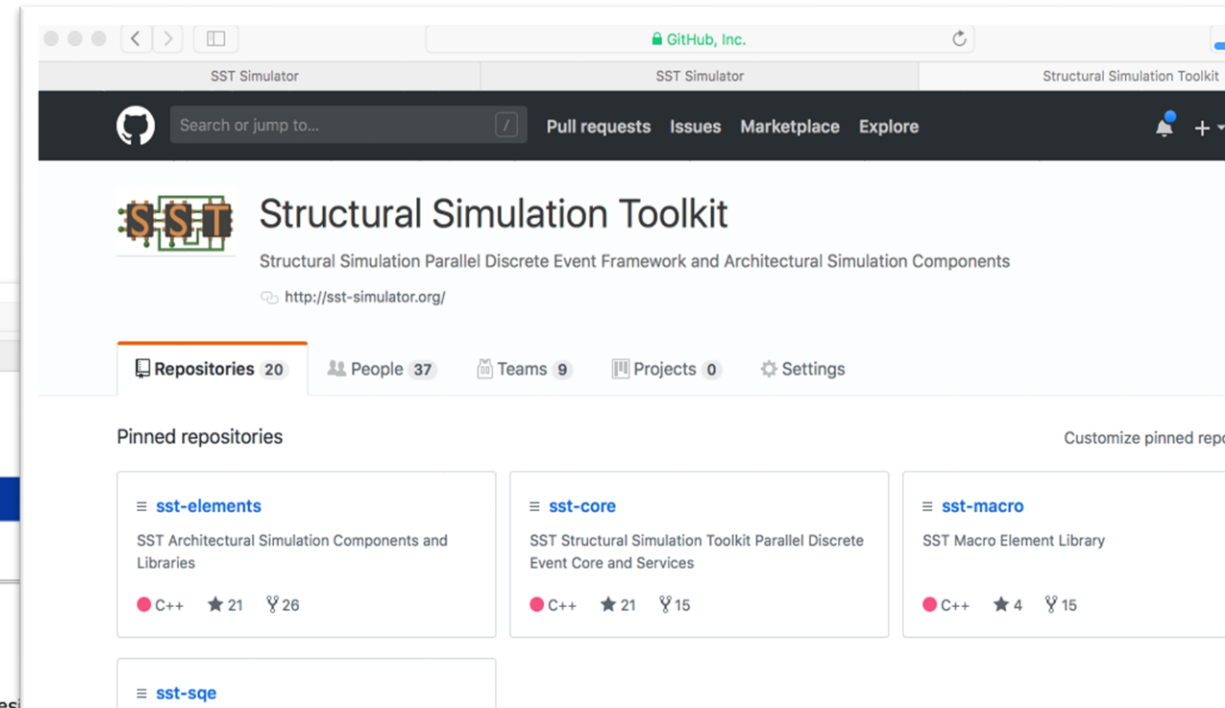
ADD CONTAINER INFO

<http://www.sst-simulator.org>

- Current release source download
- Detailed build instructions including dependencies for Linux & Mac
- Archived tutorial materials

<https://github.com/sstsimulator>

- Source code checkout
- Master branch – has passed testing
- Devel branch – has passed *basic* testing





# Building Blocks

SST simulations are comprised of **components**

Components are connected by **links**


- Every link has a minimum (non-zero) latency
- Components define **ports** which are valid connection points for a link

Components communicate by sending **events** over the links

Components can use **subcomponents** and **modules** for customizable functionality

## Element Library

*A collection of components, subcomponents, and/or modules*





# SST Code Structure



**SST Core** and **SST Elements** are compiled separately


- Element libraries *register* with the core
- External elements (not part of SST Elements) can also be registered with the core
  - Example at [github.com/sstsimulator/sst-external-element](https://github.com/sstsimulator/sst-external-element)
- Core maintains a database of registered libraries
  - Can query database with *sst-info* utility

Source code for core:

- `sst-core/src/sst/core/`

Source code for elements

- `sst-elements/src/sst/elements/`
- Most elements have a tests/ directory
  - `sst-elements/src/sst/elements/SomeComponent/tests`
  - Often a good starting point for example configurations







## Our First Simulation – demo\_1.py


We'll walk through how to configure a simulation and then run it

- Available at: <https://github.com/sstsimulator/sst-tutorials/tree/master/hpca2024/exercises>
- ( github → sstsimulator → sst-tutorials → hpca2024 )

Element libraries in our example simulation

- **Miranda** - Simple core model that runs generated instruction streams
  - Generators produce memory access patterns (SubComponent)
- **memHierarchy** – Various cache/memory system related subcomponents and modules
  - Cache (Component) with coherence protocol SubComponent
  - Memory Controller (Component) that loads a memory timing model (SubComponent)

Simulated System





## Configuration File: Global SST parameters

Set any global simulation parameters

SST Python API

User-defined string

SST argument

Other options

- Most command line options to SST are able to be set using `setProgramOption()`

```
import sst
```

```
#####
```

```
## Define SST core options
```

```
#####
```

```
# If this demo gets to 100ms, something
```

```
# has gone very wrong!
```

```
sst.setProgramOption("stop-at", "100ms")
```

Option	Definition
debug-file	File to print debug output to
heartbeat-period	If set, SST will print a heartbeat message at the specified period
print-timing-info	Tells SST to print timing information from the run
partitioner	Partitioner to use for parallel execution
output-partition	File to print partition to



## Configuration File: Declare components


SST Python API  
User-defined string  
SST argument

**Components:** `sst.Component("name", "type")`

```
#####  
## Declare components  
#####  
core = sst.Component("core", "miranda.BaseCPU")  
cache = sst.Component("cache", "memHierarchy.Cache")  
memctrl = sst.Component("memory", "memHierarchy.MemController")
```

Component name

Element library






## Configuration File: Configure the core

**Parameters:** `addParams({ "parameter" : "value", ... })`

SST Python API  
User-defined string  
SST argument

```
#####  
## Set component parameters and fill subcomponent slots  
#####  
core.addParams({  
    "clock" : "2.4GHz",  
    "max_reqs_cycle" : 2,  
})
```



*How do I know what the options are?  
Or even what elements I can pick from?*





## Aside: sst-info

Use sst-info to find information about all registered elements.

```
$ sst-info miranda.baseCPU
```

```
PROCESSED 1 .so (SST ELEMENT) FILES FOUND IN DIRECTORY(s)
```

```
[...]
```

```
=====
```

```
ELEMENT LIBRARY 0 = miranda ()
```

```
Components (1 total)
```

```
    Component 0: BaseCPU
```

```
        Description: Creates a base Miranda CPU ready to execute an address  
generator/access pattern
```

```
        ELI version: 0.9.0
```

```
        Compiled on: Dec  1 2023 14:33:14, using file: mirandaCPU.h
```

```
        Category: PROCESSOR COMPONENT
```

```
        Parameters (12 total)
```

```
            max_reqs_cycle: Maximum number of requests the CPU can issue per cycle  
            (this is for all reads and writes)  [2]
```

```
            ...
```




## Configuration File: Add a subcomponent

SST Python API  
User-defined string  
SST argument

**SubComponents:** `setSubComponent("slotname", "type")`

- Recall: SubComponent is a *swappable piece of functionality*

```
gen = core.setSubComponent("generator", "miranda.STREAMBenchGenerator")
gen.addParams({
    "n" : 1000, # Number of array elements
})
```






## Configuration File: Configure the cache

SST Python API  
User-defined string  
SST argument

```
cache.addParams({  
    "L1" : 1,  
    "cache_frequency" : "2.4GHz",  
    "access_latency_cycles" : 2,  
    "cache_size" : "2KiB",  
    "associativity" : 4,  
    "replacement_policy" : "lru",  
    "coherence_policy" : "MESI",  
    "cache_line_size" : 64,  
})
```






## Configuration File: Configure the memory controller

SST Python API  
User-defined string  
SST argument

```
memctrl.addParams({  
    "clock" : "1GHz",  
    "backing" : "none", # No real memory values, just addresses  
    "addr_range_end" : 1024*1024*1024-1,  
})  
  
memory = memctrl.setSubComponent("backend", "memHierarchy.simpleMem")  
memory.addParams({  
    "mem_size" : "1GiB",  
    "access_time" : "50ns",  
})
```






# Configuration File: Declare and connect

**Links:** `sst.Link("name")`

```
#####  
## Declare links  
#####  
core_cache = sst.Link("core_to_cache")  
cache_mem = sst.Link("cache_to_memory")  
  
#####  
## Connect components with the links  
#####  
core_cache.connect( (core, "cache_link", "100ps"),  
                    (cache, "high_network_0", "100ps") )  
  
cache_mem.connect( (cache, "low_network_0", "100ps"),  
                  (memctrl, "direct_link", "100ps") )
```

Link name





## Configuration File: Connect links


**Connect components:** `connect(endpoint1, endpoint2)`

- Where endpoint is: (component, port, latency)


```
#####  
## Connect components with the links  
#####  
core_cache.connect( (core, "cache_link", "100ps"),  
                    (cache, "high_network_0", "100ps") )  
  
cache_mem.connect( (cache, "low_network_0", "100ps"),  
                  (memctrl, "direct_link", "100ps") )
```

Endpoint 1

Endpoint 2



*How do I remember the port names?*





# SSTInfo: Getting component info

sst-info: utility to query element libraries

Optionally filter for a specific library and/or component

```
$ sst-info memHierarchy.Cache
PROCESSED 1 .so (SST ELEMENT) FILES FOUND IN DIRECTORY(s) /home/sst/build/lib/sst
Filtering output on Element = "memHierarchy.Cache"
=====
ELEMENT LIBRARY 0 = memHierarchy ()
  Component 2: Cache
    Description: Cache controller
    Parameters (41 total)
      cache_frequency: (string) Clock frequency or period with units (Hz or s; SI units OK) [<required>]
      cache_line_size: (uint) Size of a cache line (aka cache block) in bytes. [64]
    ...
    Ports (10 total)
      low_network_0: Port connected to lower level caches (closer to main memory)
    ...
    Statistics (43 total)
      TotalEventsReceived: Total number of events received, (units = "events") Enable level = 1
    ...
```

Parameter

Definition

"REQUIRED" or default value

Port name

Definition

Name

Definition

Units

Enable Level



## Running SST

Usage: `sst [options] configFile.py`

Common options:

<code>-h   --help</code>	Print complete list of command line options
<code>-v   --verbose</code>	Print information about core runtime
<code>--debug-file &lt;filename&gt;</code>	Send debugging output to specified file (default: <code>sst_output</code> )
<code>--partitioner &lt;self   simple   rrobin   linear   lib.partitioner_name&gt;</code>	Specify the partitioning mechanism for parallel runs
<code>-n   --num_threads &lt;num&gt;</code>	Specify number of threads per rank
<code>--model-options "&lt;args&gt;"</code>	Command line arguments to send to the Python configuration file. Any arguments after a final <code>-</code> will be appended to the model-options.
<code>--output-partition &lt;filename&gt;</code>	Write partitioning information to <code>&lt;filename&gt;</code>
<code>--output-dot &lt;filename&gt;</code> <code>--output-xml &lt;filename&gt;</code> <code>--output-json &lt;filename&gt;</code>	Output the configuration graph in various formats to <code>&lt;filename&gt;</code>





## Running a simulation

### Launch simulation

```
$ sst demo_1.py
```

### Output

```
Simulation is complete, simulated time: 6.80711 us
```

We probably want more information about what happened though

- Enable statistics!



## Enabling statistics

Most Components and SubComponents define statistics

```
$ sst-info memHierarchy.Cache
```

...

Statistics (43 total)

TotalEventsReceived: Total number of events received, (units = "events") Enable level = 1

CacheHits: Total number of cache hits, (units = count) Enable Level = 1

latency\_GetS\_hit: Latency for read hits, (units = cycles) Enable level = 1

### Enable statistics in the configuration file

- enableAllStatisticsForAllComponents()
- enableAllStatisticsForComponentType(type)
- enableAllStatisticsForComponentName(name)
- setStatisticLoadLevel(level)
- enableStatisticForComponentName(name, stat)
- enableStatisticForComponentType(type, stat)

### Configure output

- setStatisticOutput("sst.output\_type")
- setStatisticOutputOptions({"option" : "value", })



## Running with statistics enabled

Let's enable statistics for all components

- Caches have A LOT of statistics so send the output to a CSV file
- Other options: `sst.statoutputX` where X=
  - console
    - txt
  - json
    - hdf5

```
sst.setStatisticOutput("sst.statoutputcsv")  
  
sst.setStatisticOutputOptions({ "filepath" : "stats.csv" })  
  
sst.setStatisticLoadLevel(5)  
  
sst.enableAllStatisticsForAllComponents()
```



## Running a Simulation – Add Statistics

Copy configuration

```
$ cp demo_1.py demo_2.py
```

Add statistics to new configuration

What should you add?

Launch simulation

```
$ sst demo_2.py
```

Take a minute to look at the statistics

- Can you calculate the L1 memory bandwidth?



## SST in parallel

SST was designed from the ground up to enable scalable, parallel simulations

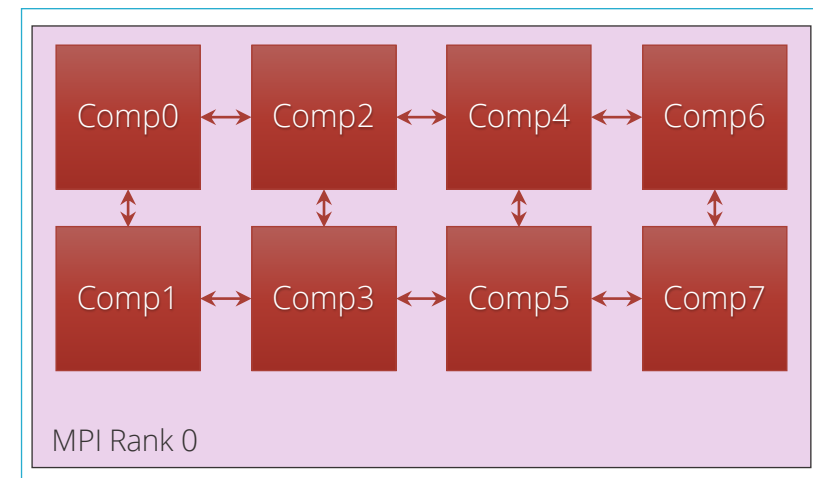
Components distributed among MPI ranks/threads

- Link latency controls synchronization rate

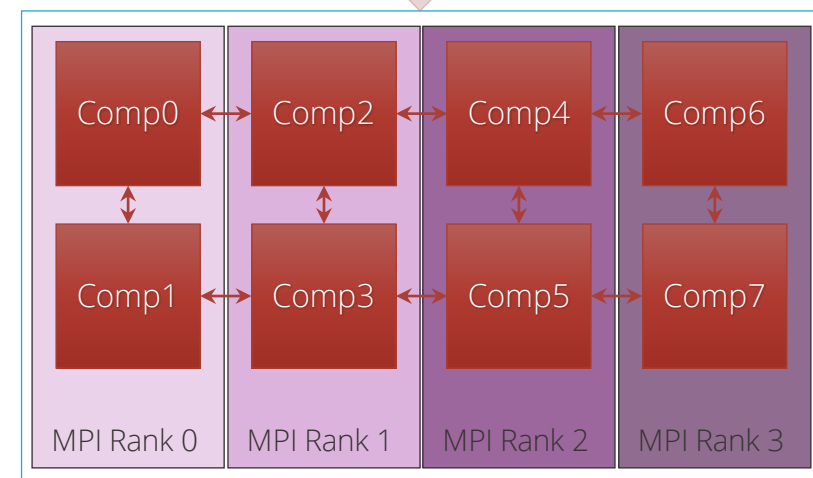
```
# Two ranks
$ mpirun -np 2 sst demo1.py

# Two threads
$ sst -n 2 demo1.py

# Two ranks with two threads each
# This will give a warning since we only
# have 3 components across 4 ranks/threads
$ mpirun -np 2 sst -n 2 demo1.py
```



Same configuration file





# A Tour of SST Elements





## SST Element Libraries

### Elements are libraries of related components

- Elements must be *registered* with the SST core
- Tells SST where to find this set of components
- Includes information on parameters and statistics for each component



### SST provides a set of element libraries

- Processor, network, memory, etc.
- Tested for interoperability within and across libraries
- Many are compatible with external “components” such as Ramulator and Spike

### You can also register your own elements

- Example: <https://github.com/sstsimulator/sst-external-element>



## (A few) SST 13.1 Elements

### Processors

- **Ariel** – PIN-based
- **Prospero** – trace execution
- **Miranda** – pattern generator
- **Vanadis** – MIPS32 and RV64 pipeline

### Memory Subsystem

- **MemHierarchy** – caches, directory, memory

### Network drivers

- **Ember** – communication patterns
- **Firefly** – communication protocols
- **Hermes** – MPI-like driver interface

### Networks/NoCs

- **Merlin** – flexible network modeling
- **Kingsley** – mesh NoC

# Processor Models

## Elements:

Ariel

Prospero

Miranda

Vanadis



## Ariel: PIN-based processor

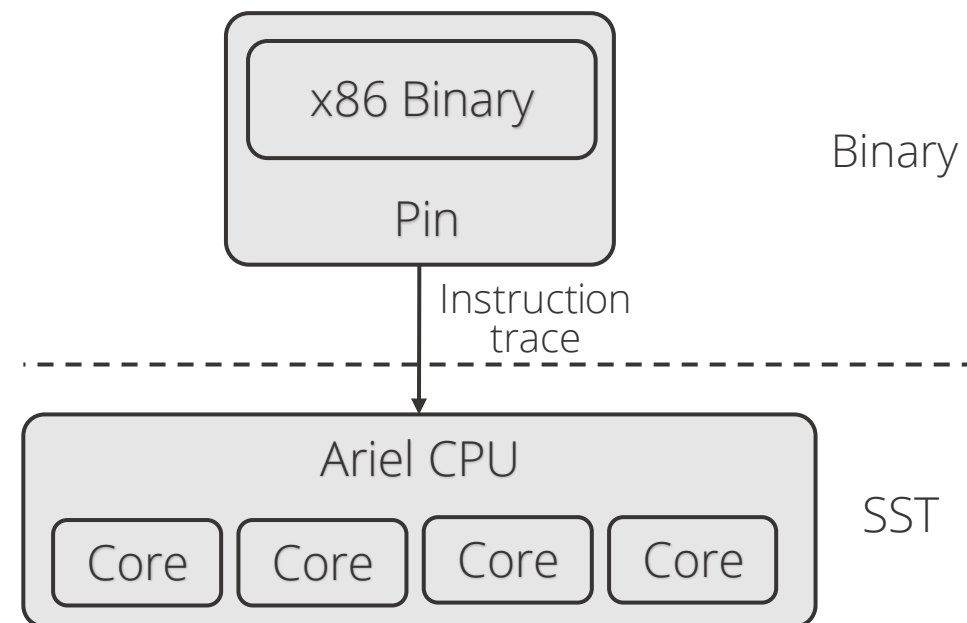
Lightweight processor core model

Uses Intel's PIN tools and XED decoders to analyze binaries

- Runs x86, x86-64, SSE/AVX, etc. binaries
- Supports fixed thread count parallelism (OpenMP, Qthreads, etc.)

Passes instructions to virtual core in SST

\$ sst-info ariel.ariel



★ GPGPU-Sim Integration



## Ariel: Details

\$ sst-info ariel.ariel

Pintool communicates with Ariel via shared memory IPC

- Per-thread FIFO of instructions from pintool to Ariel's virtual cores
- Backpressure on FIFO halts the binary's execution

Ariel's virtual cores

- **Memory instruction oriented:** execute memory instructions; other ins. single cycle no-ops
- **Clocked:** Reads instruction stream in chunks but processes on clock
- Does *not* maintain dependence order or register locations
- Can map virtual-to-physical addresses internally or use external component

Key parameters

- Ops issued/cycle
- Load/store queue size

Uses SST standardMem interface

- Generates StandardMemRequests
- Compatible with memHierarchy



## Ariel: The Tradeoff

\$ sst-info ariel.ariel

### Pros:

- Faster than more complex/pipeline models
- Reasonable approximation for studies on memory system performance
  - Especially for heavily memory-bound applications
- Reasonable model of thread interactions

### Cons

- Non-deterministic results
  - Interactions between pintool, threads, etc.
  - Variation is low ( $O(1\%)$ )
- Not compatible with non-x86 binaries
- Reliant on Pin
  - Ongoing work to enable other frontends



## Prospero: Trace-based processor

```
$ sst-info prospero.prosperoCPU
```

### Trace-based processor model

- Like Ariel, memory instruction oriented
  - Reads memory ops from a file and passes to the simulated memory system
- “Single core” but can use multiple trace files to emulate threaded or MPI applications
- Supports arbitrary length reads to account for variable vector widths
- Performs “first touch” virtual to physical mapping

### Comes with Prospero Trace Tool to generate traces

- Or can generate your own and translate to Prospero’s format



## Prospero: The Tradeoff

```
$ sst-info prospero.prosperoCPU
```

### Pros

- Faster than Ariel\*
- Provided you can get a trace

### Cons

- Traces can be very large
  - Requires good I/O system to store and read the trace
- Traces are less flexible than actual execution
  - Capture a single execution stream using a single application input





## Miranda: Pattern-based processor

\$ sst-info  
miranda.BaseCPU

Extremely light-weight processor model

- Generates memory address patterns
- Supports request dependencies

Library patterns

- Strided accesses (single stream)
  - Forward and reverse strides
- Random accesses
- GUPS
- STREAM benchmark
  - In-order & out-of-order CPU
- 3D stencil
- Sparse matrix vector multiply (SpMV)
- Copy (~array copy)
- *Stake* interface to the Spike RiscV simulator
- Ongoing work to integrate Spatter patterns



## Miranda: The tradeoffs

```
$ sst-info  
miranda.BaseCPU
```

### Pros

- Very lightweight – no binary, no trace
- Good for applications whose address patterns are predictable
  - e.g., not much pointer-chasing
- Models instruction dependences

### Cons

- Need a generator for the memory pattern of interest
  - Requires a good understanding of the pattern



## Vanadis: OOO Processor

```
$ sst-info vanadis.VanadisNodeOS
```

- MPIS32 and RISCV64 compatible processor model
- OOO model
  - Configure micro-architectural details
    - Instruction fetch/decode/retire rate
    - ROB size
  - Branch prediction
- Multi-threaded cores
- Musl libc used to cross-compile programs
- System-call emulation



## Vanadis: The tradeoffs

```
$ sst-info vanadis.VanadisNodeOS
```

### Pros

- Runs binaries
- Detailed model of instruction dependencies

### Cons

- Slower than other models

# Memory

Elements:  
memHierarchy



## MemHierarchy: Memory system

```
$ sst-info memHierarchy | grep "Component"
```

Collection of interoperable memory system elements

- Caches
- Directories
- Memory controllers
  - Interfaces to memory models (DDR, HBM, HMC, NVM, etc.)
- Scratchpads
- NoC (network-on-chip) interfaces
- Buses

Components are cycle-accurate/cycle-level

Capable of modeling modern cache and memory subsystems



## MemHierarchy: Cache modeling

\$ sst-info memHierarchy.Cache

### Highly configurable

- Arbitrary hierarchy depth, flexible topologies
- Cache inclusivity, coherence, private/shared, etc. configurable
- Single- and multi-socket configurations
- Prefetch via *Cassini* element library

### Data movement

- Components support direct, bus, and on-chip network (NoC) communication

Event types: read/write, atomics, LLSC, noncacheable, custom memory, etc.



## MemHierarchy: Memory modeling

```
$ sst-info memHierarchy.MemController
```

Interface to memory is the *MemController*

MemControllers implement *backends*

- These do the actual work of timing memory access
- Can be interfaces to other memory simulators
- More on the next slide

Support *custom memory instructions*

- Including ability to do cache shutdowns for coherence maintenance





## MemHierarchy: Memory modeling


\$ sst-info memHierarchy.MemController

### Memory controller

- Manages data values if needed (backing store)
- Facilitates custom memory commands
  - Including cache shutdowns for coherence maintenance
- Passes events to *memory backend* subcomponent

### Backend: the “real” memory controller and/or memory

- Implementations
  - Memory controller and model itself
  - Memory controller with interface to a memory component
  - Interface to another memory controller/memory component
  - Wrapper to an external simulator





## MemHierarchy: SST 13.1 backends

### Memory (external)

- CramSim (DDR, HBM)
  - **DRAMSim3** (DDR, LPDDR, GDDR HBM, HMC, STT-MRAM)
  - **FlashDIMMSim** (FLASH)
  - **HMCSim/GoblinHMC** (HMC)
  - Messier (NVRAM)
  - **Ramulator** (DDR, HBM, HMC)
  - SimpleDRAM (DDR)
  - SimpleMem (constant latency)
  - TimingDRAM (DDR)
  - VaultSimC (HMC-like)
- Plus a few that can be used with other backends to reorder requests, add latency, etc.



## Running a Simulation – Add Components, L2 Cache

Copy configuration

```
$ cp demo_2.py demo_3.py
```

Add an L2 cache between L1 and memory to new configuration

What should you add?

- What parameters are available for an L2 cache?
- What are appropriate values for the parameters?

Launch simulation

```
$ sst demo_3.py
```

- How did this affect your overall simulated time?
- How did this affect traffic to and from your backing store?



## Running a Simulation – Switch Components, Timing DRAM

Copy configuration

```
$ cp demo_3.py demo_4.py
```

Switch the simpleMem subcomponent for timingDRAM

What should you change? Remember that sst-info is your pal!

Launch simulation

```
$ sst demo_4.py
```

- How do your results differ from the run with simpleMem?

# Networks

Elements:

Merlin  
Kingsley



## Merlin: Network simulator

\$ sst-info merlin

Low-level networking components that can be used to simulate high-speed networks (machine level) or on-chip networks

### Capabilities

- High radix router model (hr\_router)
- Topologies – mesh, n-dim tori, fat-tree, dragonfly,

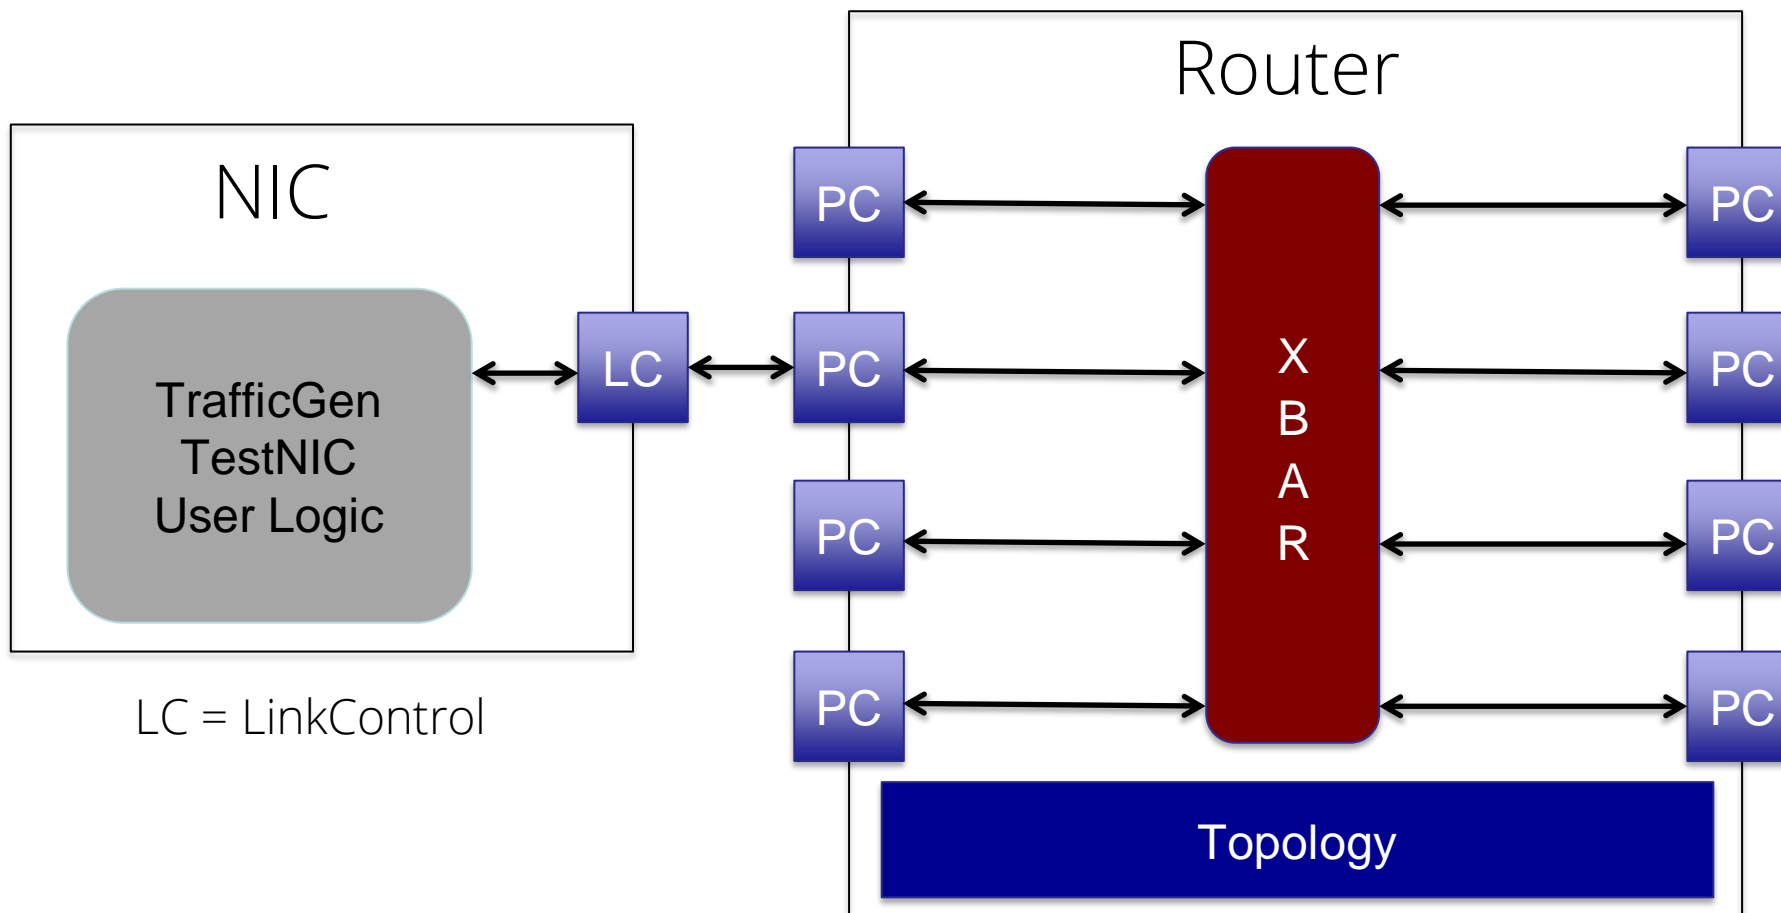
### Many ways to drive a network

- Simple traffic generation models
  - Nearest neighbor, uniform, uniform w/ hotspot, normal, binomial
- *MemHierarchy*
- Lightweight network endpoint models (*Ember* – coming up next)
- Or, make your own



# Merlin: Organization

\$ sst-info merlin



LC = LinkControl

PC = PortControl




## Kingsley: Mesh simulator

\$ sst-info kingsley

Network-on-chip model; mesh configuration

Similar to Merlin but:

- No input queuing at routers
- Mesh topology only
- Not all ports need to be populated
- Possible to instantiate multiple unconnected networks
- Multiple physical networks for coherence (e.g., request/response/ack/forward)
- Kingsley NoC + Merlin/Kingsley system network





# Network Drivers

Elements:

Ember

Firefly

Hermes



## Ember: Network Traffic Generator

\$ sst-info ember

Light-weight endpoint for modeling network traffic

- Enables large-scale simulation of networks where detailed modeling of endpoints would be expensive

Packages patterns as *motifs*

- Can encode a high level of complexity in the patterns
- Generic method for users to extend SST with additional communication patterns

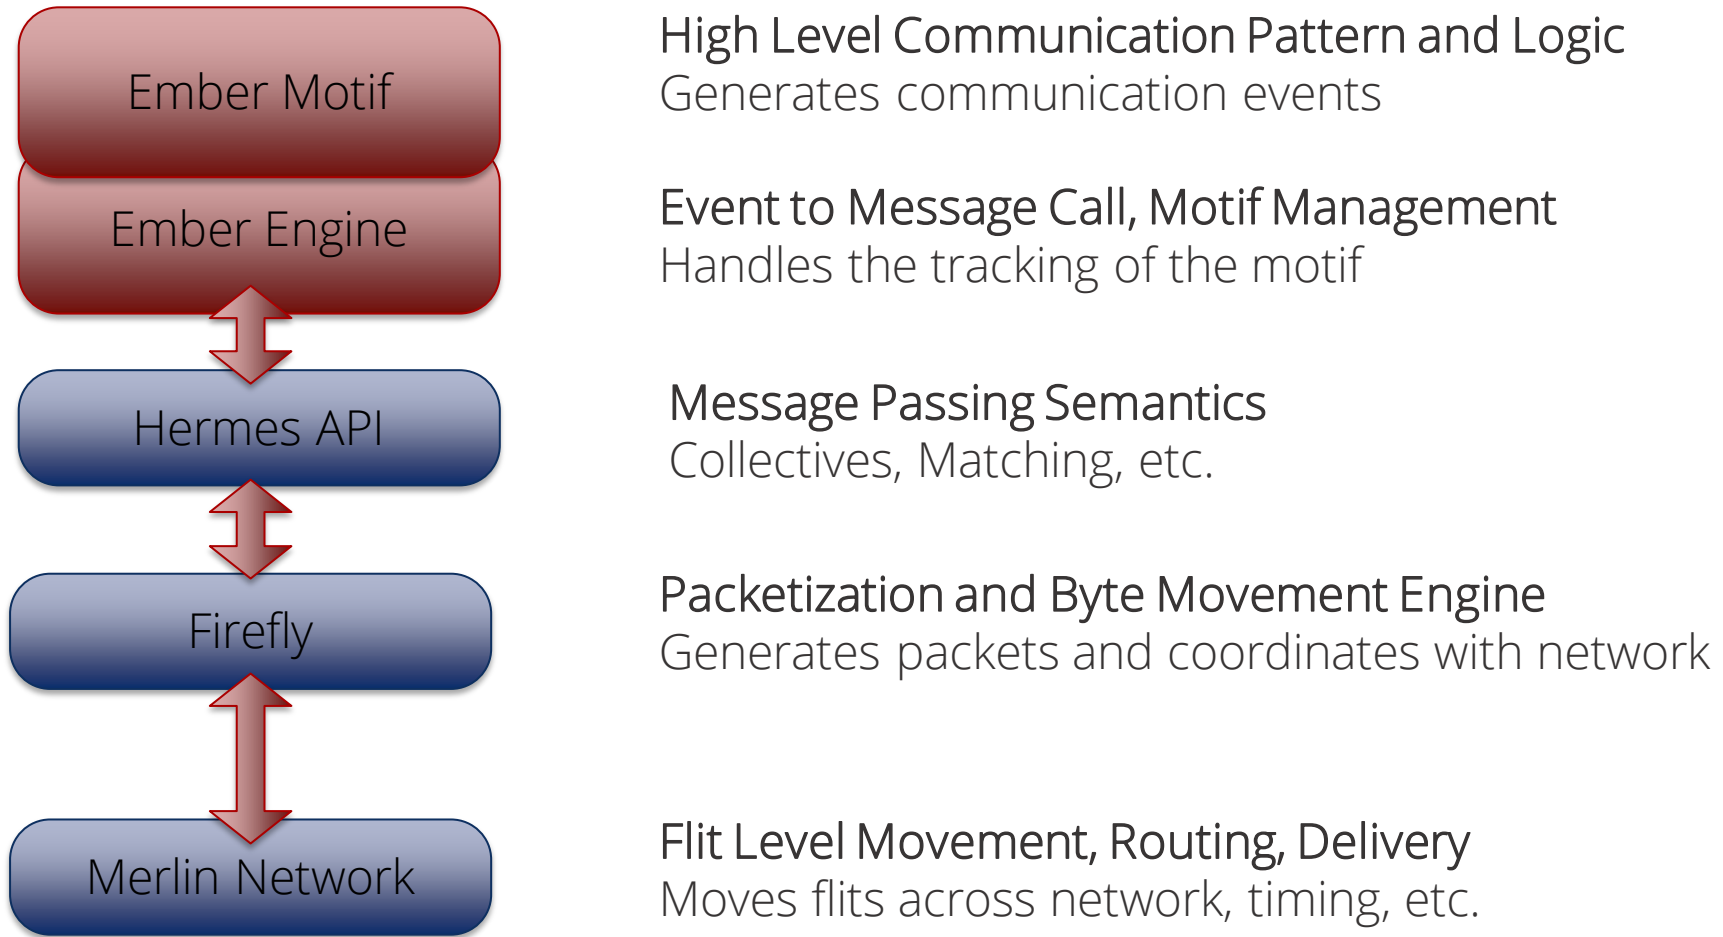
Intended to be a driver for the Hermes, Firefly, and Merlin communication modeling stack

- Uses Hermes message API to create communications
- Abstracted from low-level, allowing modular reuse of additional hardware models



## Ember: Overview

\$ sst-info ember





## Ember: Motifs

```
$ sst-info ember | grep Motif
```

Motifs are lightweight patterns of communication

- Tend to have very small state
- Extracted from parent applications
- Models as an MPI program (serial flow of control)
  - Many motifs acting in the simulation create the parallel behavior

### Example motifs

- Halo exchanges (1, 2, and 3D)
- MPI collections – reductions, all-reduce, gather, barrier
- Communication sweeping (Sweep3D, LU, etc.)



## Ember: Motifs (continued)

```
$ sst-info ember.EmberEngine
```

The EmberEngine creates and manages the motif

- Creates an event queue which the motif adds events to when probed
- The Engine executes the queued events in order, converting them to message semantic calls as needed
- When the queue is empty, the motif is probed again for events

Events correspond to a specific action

- E.g., send, recv, allreduce, compute-for-a-period, wait, etc.



## Firefly: Network traffic

\$ sst-info firefly

**Purpose:** Create network traffic, based on application communication patterns, at large scale

- Enables testing the impact of network topologies and technologies on application communication at very large scale

Scales to 1 million nodes

Supports multiple “cores” per Node

- Interaction between cores limited to message passing

Supports space sharing of the network

- Multiple “apps” running simultaneously



## Firefly: Simulating large networks

\$ sst-info firefly

A network node consists of

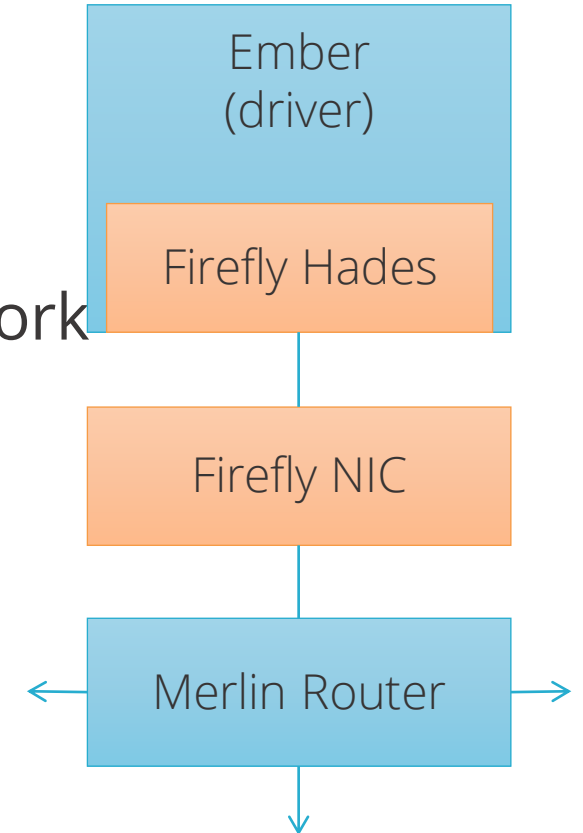
- Driver (the “application”)
- NIC
- Router

Nodes are connected together via routers to form a network

- Fat tree, torus, etc.

Firefly is the interface between the driver and the router

- Message passing library → Firefly Hades
- NIC → Firefly NIC





## (More) SST 13.1 Elements

### Processors

- **Ariel** – PIN-based
- **Prospero** – trace execution
- **Miranda** – pattern generator
- **Vanadis** – MIPS32 and RV64 pipeline

### Memory Subsystem

- **MemHierarchy** – caches, directory, memory

### Network drivers

- **Ember** – communication patterns
- **Firefly** – communication protocols
- **Hermes** – MPI-like driver interface

### Networks/NoCs

- **Merlin** – flexible network modeling
- **Kingsley** – mesh NoC





## Even More Elements

There are even more elements included, described here:

- <http://sst-simulator.org/sst-docs/docs/elements/intro>

A number of external simulators are compatible with SST.

- See section on optional dependencies: [http://sst-simulator.org/SSTPages/SSTBuildAndInstall\\_13dot1dot0\\_SeriesDetailedBuildInstructions/](http://sst-simulator.org/SSTPages/SSTBuildAndInstall_13dot1dot0_SeriesDetailedBuildInstructions/)
- Many memory timing simulators are supported by memHierarchy



## Viewing Configuration Graph

Let's take a look at how SST views our system

Re-run demo\_4 but add a command to dump the configuration graph

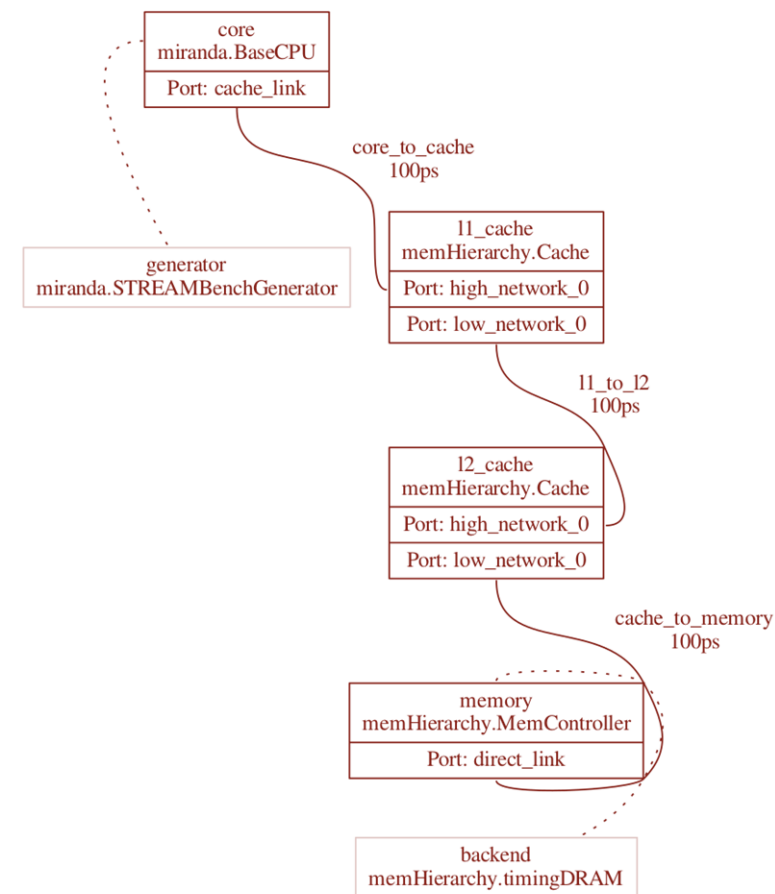
```
$ sst --output-dot=graph_demo_4.dot demo_4.py
```

This gives you a GraphViz formatted file

```
$ dot -Tpdf graph_demo_4.dot -o graph_demo_4.pdf  
$ evince graph_demo_4.pdf
```

Is this how you expected your system to look?

With this in mind, let's add a second Miranda core...





## Running a Simulation – Add Components, Second Miranda

Copy configuration


```
$ cp demo_4.py demo_5.py
```

Add an a second Miranda generator

- What else might you need?
- How about another L1 cache?
- How are you going to wire everything together? How about a bus?

Dump the wiring diagram to verify the model  
Launch simulation

```
$ sst demo_5.py
```





## Running a Simulation – Add Components, Second L2


Copy configuration

```
$ cp demo_5.py demo_6.py
```

Add L2 cache and move both above the bus

Launch simulation

```
$ sst demo_6.py
```





# Running a Simulation – Add Components, Secondary Memory

## Copy configuration

```
$ cp demo_6.py demo_7.py
```

## Add an a second L2 and memory controller


- Think carefully about how addressing should will work...

Can you still use a bus?

Can the talk directly to the memory controller?

## Launch simulation

```
$ sst demo_7.py
```





# Getting Help & Extending SST



## Extending SST

SST was designed for extensibility

- Components/subcomponents can be added without touching SST Elements
  - Example: write a new prefetcher and have memH caches use it → *no changes* to memHierarchy
- SST-Core APIs are stable → one year deprecation period
  - Element APIs may be less so but generally try to keep them consistent
- Many users start with SST Elements and then build their own customized libraries
  - Partially or completely replacing SST Element functionality

Many approaches to using SST

- Core only: Write your own components from scratch
- Start from existing Elements and replace components/subcomponents to meet your needs
- Wrap existing simulators and insert as components or subcomponents



## Extending SST: Resources

### Example element library

- Components demonstrating links, ports, clocks, event handling, etc.
- `sst-elements/src/sst/elements/simpleElementExample/`

### simpleSimulation

- Simulates a car wash (a little more complex than example elements)

### Example external element library

- Demonstrates building and registering a new element library
- <https://github.com/sstsimulator/sst-external-element>

### Website

- *Getting Started Extending SST (a little out of date)*
- *Building Element Libraries outside SST source tree*
- Past tutorial material (under Downloads)
- [sst-simulator.org/SST-website](https://sst-simulator.org/SST-website) API documentation





## Finally: Getting help

SST website contains lots of information ([www.sst-simulator.org](http://www.sst-simulator.org))

- Downloading, installing, and running SST
- Element libraries and external components
- Guides for extending SST
- Information on APIs
- Information about current development efforts
- Past tutorial slides and exercises

### SST Github

- Current development
- Issues track user questions as well as development plans, bugs, etc.



## Part 1 wrap-up

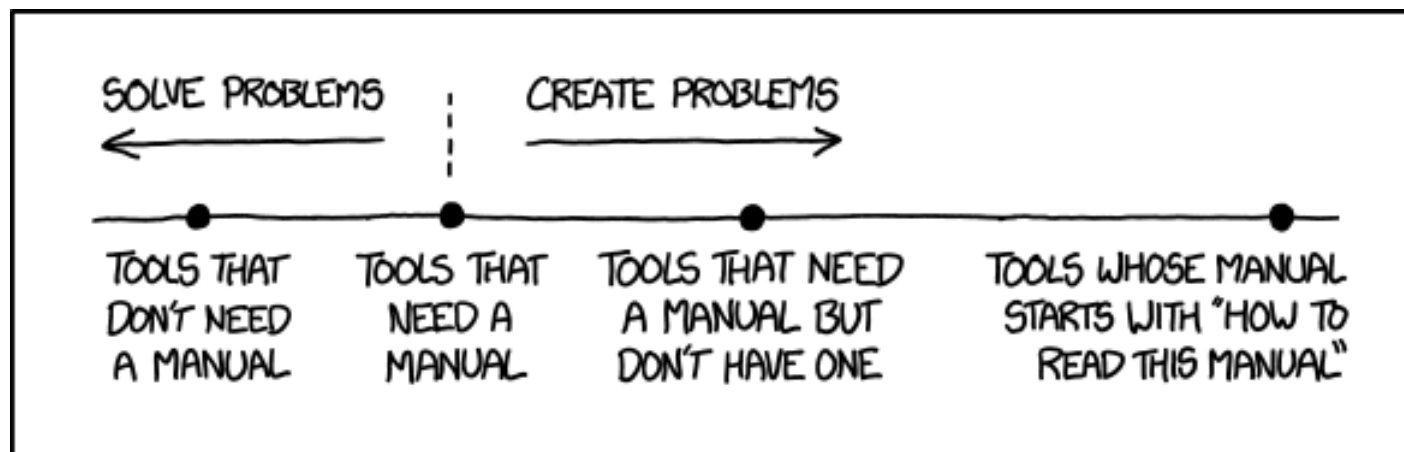
SST is a parallel, flexible simulation framework

- Can simulate many systems at many granularities
- Capable of simulating modern architectures
- Modular design for extensibility

Please keep us posted on your uses of SST as well as any capabilities you've added or would like to see added

The SST team wants to help you!

- Documentation?
- Examples?
- Kittens?



A dark blue notebook with a textured cover is shown in the top left corner. Two pens, one silver and one black, are resting on the notebook. A white cup of coffee with a dark liquid inside is in the bottom left corner. A diagonal line with a multi-colored pattern (green, yellow, orange, red, purple, blue) runs from the bottom left towards the top right, separating the notebook area from the rest of the slide.

**BREAK**

Return at 10:20



# Advanced Node Models



## Welcome!

Part 1: Introduction to SST

8:00 – 10:00

Break

10:00 – 10:20

### ***Part 2: Node-Level Simulation***

***10:20 – 11:00***

Part 3: System-Scale Simulation

11:00 – 12:20

SST Overview


Hands on: Basic Simulation in SST

A Tour of SST Elements

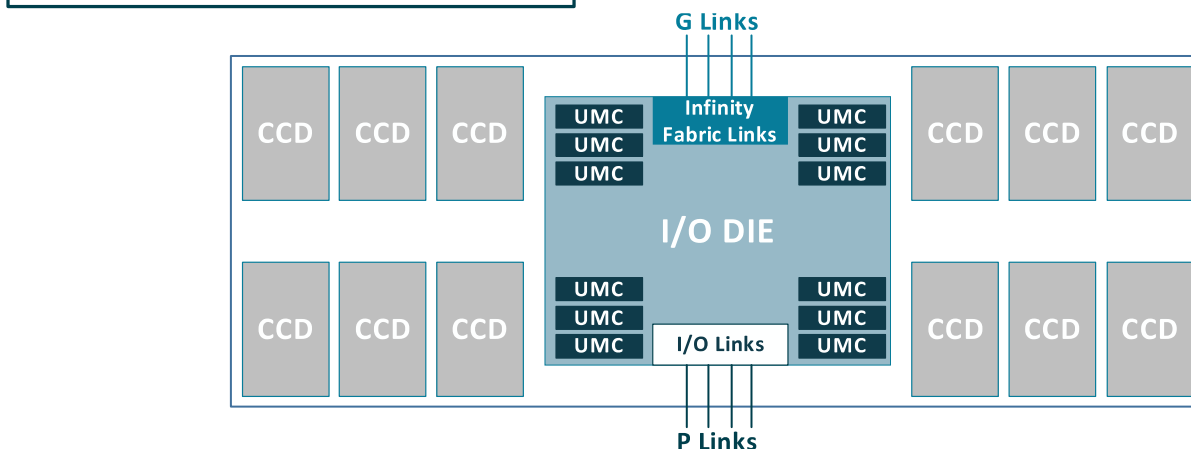



# SOCs

What if you wanted something much more complex? Can SST simulate a core with SMT? What about something like a modern processor? SPR? EPYC?



<https://www.intel.com/content/www/us/en/developer/articles/technical/fourth-generation-xeon-scalable-family-overview.html>



<https://www.amd.com/content/dam/amd/en/documents/epyc-technical-docs/white-papers/58015-epyc-9004-tg-architecture-overview.pdf>






## Single-core Vanadis

Let's start simple...

- RISC-V core
- I-Cache
- D-Cache
- L2 Cache
- Memory





# Single-core Vanadis

What attributes can we change in the core model?

```
$ sst-info vanadis
```

## Parameters (35 total)

verbose: Set the level of output verbosity, 0 is no output, higher is more output [0]  
dbg\_mask: Mask for output. Default is to not mask anything out (0) and defer to 'verbose'. [0]  
start\_verbose\_when\_issue\_address: Set verbose to 0 until the specified instruction address is issued, then set to 'verbose' parameter []  
stop\_verbose\_when\_retire\_address: When the specified instruction address is retired, set verbose to 0 []  
pause\_when\_retire\_address: If specified, the simulation will stop when this address is retired. [0]  
pipeline\_trace\_file: If specified, a trace of the pipeline activity will be generated to this file. []  
max\_cycle: Maximum number of cycles to execute. The core will halt after this many cycles. [std::numeric\_limits<uint64\_t>::max()]  
node\_id: Identifier for the node this core belongs to. Each node in the system needs a unique ID between 0 and (number of nodes) - 1. Used to tag output. [0]  
core\_id: Identifier for this core. Each core in the system needs a unique ID between 0 and (number of cores) - 1. [<required>]  
hardware\_threads: Number of hardware threads in this core [1]  
clock: Core clock frequency [1GHz]  
reorder\_slots: Number of slots in the reorder buffer [64]  
physical\_integer\_registers: Number of physical integer registers per hardware thread [128]  
physical\_fp\_registers: Number of physical floating point registers per hardware thread [128]  
integer\_arith\_units: Number of integer arithmetic units [2]  
integer\_arith\_cycles: Cycles per instruction for integer arithmetic [2]  
integer\_div\_units: Number of integer division units [1]  
integer\_div\_cycles: Cycles per instruction for integer division [4]  
fp\_arith\_units: Number of floating point arithmetic units [2]  
fp\_arith\_cycles: Cycles per floating point arithmetic [8]  
fp\_div\_units: Number of floating point division units [1]  
fp\_div\_cycles: Cycles per floating point division [80]  
branch\_units: Number of branch units [1]  
branch\_unit\_cycles: Cycles per branch [int\_arith\_cycles]  
issues\_per\_cycle: Number of instruction issues per cycle [2]  
fetches\_per\_cycle: Number of instruction fetches per cycle [2]  
retires\_per\_cycle: Number of instruction retires per cycle [2]  
decodes\_per\_cycle: Number of instruction decodes per cycle [2]  
dcache\_line\_width: Width of a line for the data cache, in bytes. (Currently not used but may be in the future). [64]  
icache\_line\_width: Width of a line for the instruction cache, in bytes [64]  
print\_retire\_tables: Print registers during retirement step (default is yes) [true]  
print\_issue\_tables: Print registers during issue step (default is yes) [true]  
print\_int\_reg: Print integer registers true/false, auto set to true if verbose > 16 [false]  
print\_fp\_reg: Print floating-point registers true/false, auto set to true if verbose > 16 [false]  
print\_rob: Print reorder buffer state during issue and retire [true]





# Single-core Vanadis

What attributes can we change in the core model?

```
$ sst-info vanadis
```

Default Parameter

Parameters (Purged a Few...)

```
clock: Core clock frequency [1GHz]
reorder_slots: Number of slots in the reorder buffer [64]
physical_integer_registers: Number of physical integer registers per hardware thread [128]
physical_fp_registers: Number of physical floating point registers per hardware thread [128]
integer_arith_units: Number of integer arithmetic units [2]
integer_arith_cycles: Cycles per instruction for integer arithmetic [2]
integer_div_units: Number of integer division units [1]
integer_div_cycles: Cycles per instruction for integer division [4]
fp_arith_units: Number of floating point arithmetic units [2]
fp_arith_cycles: Cycles per floating point arithmetic [8]
fp_div_units: Number of floating point division units [1]
fp_div_cycles: Cycles per floating point division [80]
branch_units: Number of branch units [1]
branch_unit_cycles: Cycles per branch [int_arith_cycles]
issues_per_cycle: Number of instruction issues per cycle [2]
fetches_per_cycle: Number of instruction fetches per cycle [2]
retires_per_cycle: Number of instruction retires per cycle [2]
decodes_per_cycle: Number of instruction decodes per cycle [2]
```



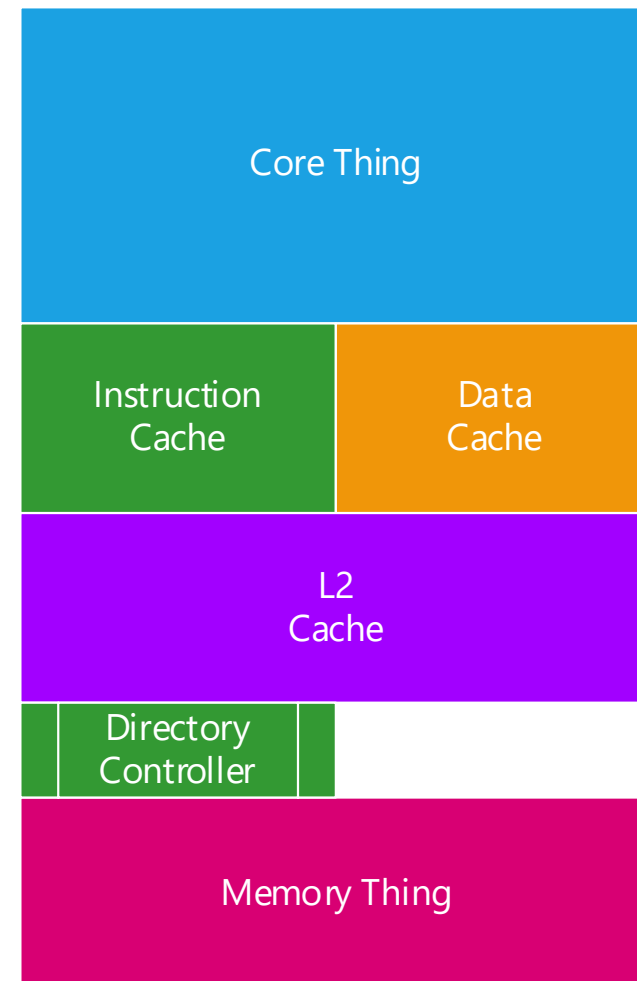
# Single-core Vanadis

Let's start simple...

```
$ cat no_rtr_vanadis.py
```

Ok, not so simple but the structure is similar to all SST configurations

```
$ sst no_rtr_vanadis.py
```






## Multi-core Vanadis

What if we want to make something more complex?

- 3x2 mesh
- CPUs at 1, 3, 4, 5
- Memory at 0, 2






## Multi-core Vanadis

What if we want to make something more complex?

```
$ cat boom_vanadis-kingsley.py
```

This looks frightening but isn't that bad...

```
$ sst boom_vanadis-kingsley.py
```



Thanks!

Exceptional Service In The National Interest



## Join Us!

We are always looking for new staff and new collaborations...

- Design challenges in the post-exascale era requires that we draw from a diverse pool of talent across multiple disciplines!

If you're interested, check out <https://www.sandia.gov/careers/career-possibilities/career-opportunities/computer-science/>, or contact Clay at [chughes@sandia.gov](mailto:chughes@sandia.gov)



**Backup**





```
$ cp demo_6.py demo_7.py
```

- Each PE has a Miranda generator and an L1
- The two memory nodes should use timingDRAM
- How should you connect everything?

```
$ sst demo_7.py
```

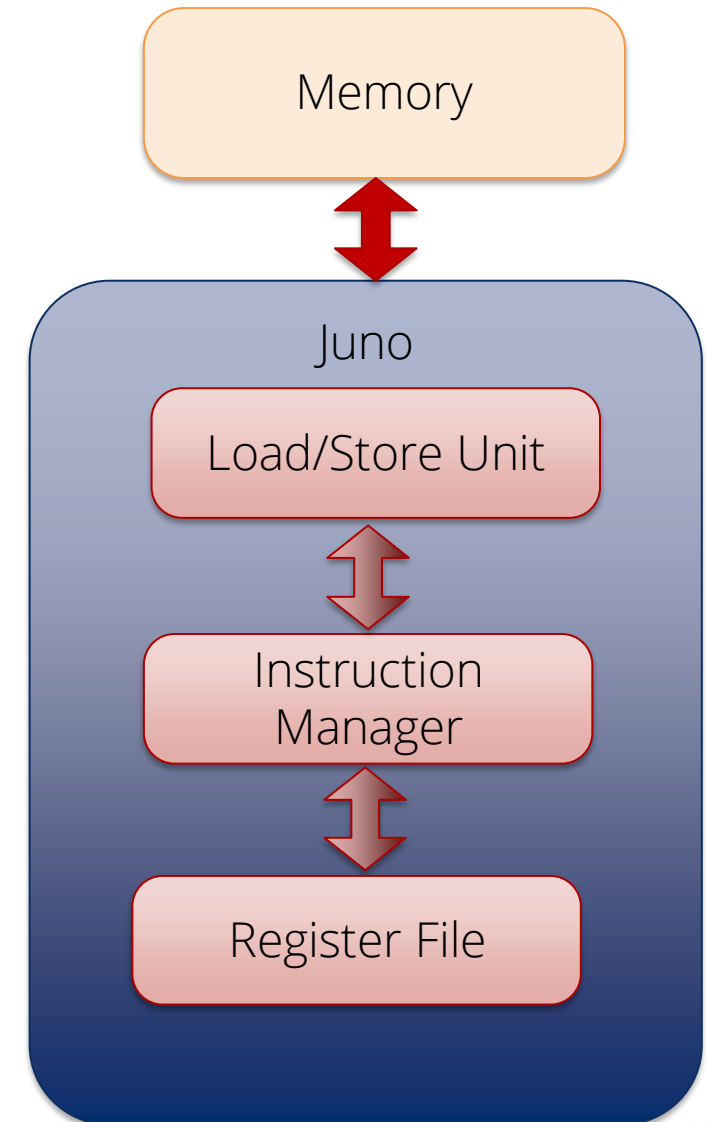




## Juno: Simple instruction processor

Executes a program written in simple “assembly”

- 32-bit wide instructions with 8 bit op codes
- 64-bit integer operations
  - ADD, SUB, DIV, MUL,
  - AND, OR, XOR, NOT
- Jump by register value (JGT-Zero, JLZ-Zero, J-Zero)
  - Jump up to 16 bits in either direction from current PC
- Up to 253 user registers
  - r0 = PC
  - r1 = data start register





# Running a Simulation – Add Components, Second Memory

## Copy configuration

```
$ cp demo_6.py demo_7.py
```

## Add an a second L2 and memory controller


- Think carefully about how addressing should will work...

Can you still use a bus?

Can the talk directly to the memory controller?

## Launch simulation

```
$ sst demo_7.py
```




# Running a Simulation – Add Components, Second Memory

Swap the bus component for the shogun component

```
shogun_xbar = sst.Component("shogunxbar",
    "shogun.ShogunXBar")
shogun_xbar.addParams({
    "clock" : "1.0GHz",
    "port_count" : 4,
    "verbose" : 0
})
```

Add directory controllers before the memory

```
dirctrl = sst.Component("dirctrl_" + str(cache_id), "memHierarchy.DirectoryController")
dirctrl.addParams({
    "coherence_protocol" : "MESI",
    "entry_cache_size" : "32768",
    "addr_range_end" : endAddr,
    "addr_range_start" : startAddr,
    "interleave_size" : "256B",
    "interleave_step" : str(numLLC * 256) + "B",
})
dc_cpulink = dirctrl.setSubComponent("cpulink", "memHierarchy.MemNIC")
dc_memlink = dirctrl.setSubComponent("memlink", "memHierarchy.MemLink")
dc_cpulink.addParams({
    "group" : 3,
})
dc_linkctrl = dc_cpulink.setSubComponent("linkcontrol", "shogun.ShogunNIC")
```





Slide Left Blank